

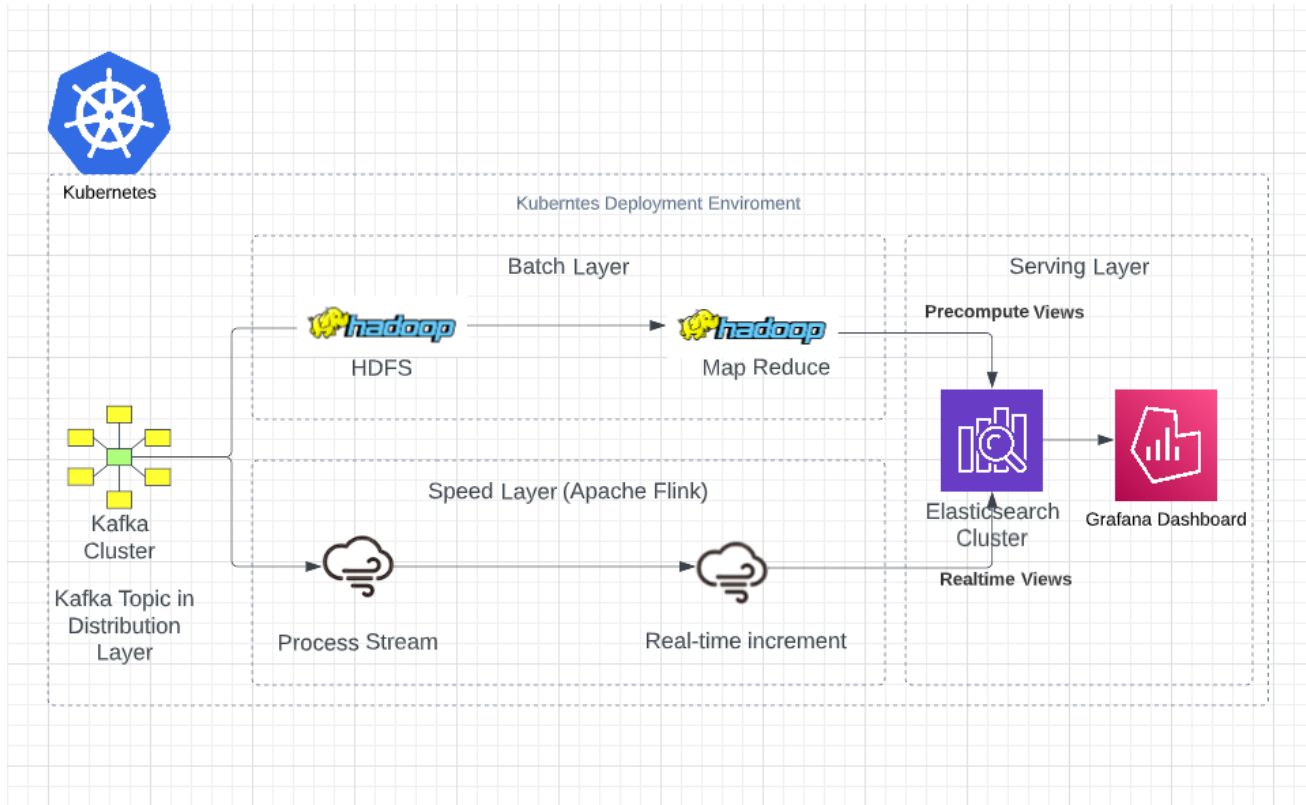
**CMM705**  
**COURSEWORK**

**Name:** G. A. J. M Senadeera

**ID:** 2121903 (RGU) / 20211475(IIT)

## 1. Solution Architecture

1. This system should have the ability to analyze real-time and historical data. So, I would recommend implementing the most prominent Lambda architecture because it supports both historical and real time data views to end users to query.



2. This proposed system is planned to deploy in a **Kubernetes** Deployment environment because **CNCF** supports each of these items to deploy in such environments ( <https://landscape.cncf.io/> ). And let's describe each section/component in this architecture.

### Distribution Layer:

This layer mainly consisted of a **Kafka cluster**, this messaging broker is widely used in the industry because of multiple advantages of using it (performance, deployment support, features). In the broker, in this design, we are interested in a **single topic** which will be receiving real time cricket data. This data will be **multiplexed** into the speed and batch layers for processing.

### Batch Layer:

This layer is considering processing legacy data that will be **gathered and accumulated** over time from the **Kafka topic** which we discussed earlier, and we call the data as **master dataset**, which is **immutable and incremental**. **Hadoop HDFS cluster** is used to save this dataset. And then **MapReduce** jobs will be triggered to processing on the dataset

periodically. HDFS itself provides reliable and durable distributed datastore and it could manage to scale out over higher volumes of data received over the time.

### **Speed Layer:**

This layer is mainly focused on processing real time data that has not been replicated in the **master dataset**. This incremental portion will be processed in real time using **Apache Flink** and saved in incremental views which is a new technology compared to **Apache Strom**, it provides high throughput and seamless integration with data sinks with simpler coding.

### **Serving Layer:**

So far, we have discussed about two processing layers, but we didn't discuss where to point out the data after processing them. After each process/layer that we have discussed, the data will be saved in an **Elasticsearch**, NoSQL database. **Flink** natively supports **Elasticsearch** as a data **sink**, while we can configure **Hadoop MapReduce** Jobs to save processed data in the **Elasticsearch** by configuring **elasticsearch-hadoop.jar** and implementing the code which is provided by the **Elasticsearch team** to integrate with **Hadoop**.

Advantage of using this is, we will have single source of truth of the whole data that has been processed and further integration will also be simple with the other services because of this fact.

In the serving layer, we have also a **Grafana** deployment which will be looking at the **Elasticsearch** database and query data from it into its dashboard components. The reason for using this combination of **Elasticsearch and Grafana** is, they have seamless integration with each other. On the other hand, **Grafana** is widely used in the industry because it has mature features, and **Elasticsearch** is distributed analytics engine for free.

## 2. Data Analysis

### 2.1. Analyze the following using Hadoop MapReduce

For the analysis, I use **Kubernetes** cluster to deploy **hadoop-hive-pig** container because I don't have enough resources to run this container on my local machine. The following prerequisite steps would be the same for both (2.1, 2.2) questions.

- i. First, switch the context to get into the correct cluster.

**kubectl config use-context [se-prod-k8s-user@se.prod.cluster.local](#)**

```
jmadushan@CLLK-JANITHAM:~$ kubectl config use-context se-prod-k8s-user@se.prod.cluster.local
Switched to context "se-prod-k8s-user@se.prod.cluster.local".
jmadushan@CLLK-JANITHAM:~$
```

- ii. Create a pod using the docker image **suhothayan/hadoop-hive-pig:2.7.1**

**kubectl run hdfs --image suhothayan/hadoop-hive-pig:2.7.1 --restart Never**

```
jmadushan@CLLK-JANITHAM:~$ kubectl run hdfs --image suhothayan/hadoop-hive-pig:2.7.1 --restart Never
pod/hdfs created
jmadushan@CLLK-JANITHAM:~$
```

- iii. Now copy the given **dataset** into the **pod/container** that has been created by the previous step.

**kubectl cp ipl/ipl-data.csv hdfs:/**

Let's see the copied content

**kubectl exec -it hdfs -- ls**

```
jmadushan@CLLK-JANITHAM:~/rgu/CMM705$ kubectl exec -it hdfs -- ls
bin          dev          ipl-data.csv  media        opt          sbin         sys         var
boot         etc          lib           metastore_db proc          selinux      tmp
derby.log    home        lib64         mnt          root         srv          usr
jmadushan@CLLK-JANITHAM:~/rgu/CMM705$
```

- iv. Connect into the pod that we have created

**kubectl exec -it hdfs -- sh**

```
jmadushan@CLLK-JANITHAM:~/rgu/CMM705$ kubectl exec -it hdfs -- sh
sh-4.1#
```

Create directory called **ipl** in the **hdfs**

**hdfs dfs -mkdir /ipl**

```
jmadushan@CLLK-JANITHAM:~/rgu/CMM705$ kubectl exec -it hdfs -- sh
sh-4.1# hdfs dfs -mkdir /ipl
22/12/05 11:57:05 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
sh-4.1#
```

And put the dataset **ipl-data.csv** into the folder that we have created in the **hdfs**

**hdfs dfs -put ipl-data.csv /ipl/ipl-data.csv**

```
22/12/05 11:57:05 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
sh-4.1# hdfs dfs -put ipl-data.csv /ipl/ipl-data.csv
22/12/05 11:59:17 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
sh-4.1#
```

MapReduce code has been included in the folder called **coursework-mapreduce-code**.

In order to run the map-reduce jobs, we need to build the related jar first using following command.  
**mvn clean package**

```
[INFO] No sources to compile
[INFO]
[INFO] --- maven-surefire-plugin:2.12.4:test (default-test) @ coursework-mapreduce ---
[INFO] No tests to run.
[INFO]
[INFO] --- maven-jar-plugin:2.4:jar (default-jar) @ coursework-mapreduce ---
[INFO] Building jar: C:\workspace\iit\cmm705\cw\course-work-mapreduce\target\coursework-mapreduce-1.0-SNAPSHOT.jar
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 10.314 s
[INFO] Finished at: 2022-12-06T07:17:50+05:30
[INFO] -----

Process finished with exit code 0
|
```

And copy the jar file (**coursework-mapreduce-1.0-SNAPSHOT.jar**) into the **hdfs** container that we have already created.

**kubect1 cp target/coursework-mapreduce-1.0-SNAPSHOT.jar hdfs:/**

```
madushan@CLLK-JANITHAM:~/ws/iit/cmm705/cw/course-work-mapreduce$ kubect1 cp target/
lasses/                               generated-sources/           maven-status/
coursework-mapreduce-1.0-SNAPSHOT.jar  maven-archiver/
madushan@CLLK-JANITHAM:~/ws/iit/cmm705/cw/course-work-mapreduce$ kubect1 cp target/coursework-mapreduce-1.0-SNAPSHOT.jar hdfs:/
madushan@CLLK-JANITHAM:~/ws/iit/cmm705/cw/course-work-mapreduce$
```

The reduction implementation that I am going to use for both questions, is the same as the following mentioned. It will count the number of occurrences of each key and save write it back to the HDFS.

*The implementation is mostly the same that we have used in the lab sessions in the class.*

```
public class IntSumReducer extends
    Reducer<Text, IntWritable, Text, IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(
        final Text key,
        final Iterable<IntWritable> values,
        final Context context
    ) throws IOException, InterruptedException {
        // Count the occurrences with the same key
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        // Set the count to the result object
        result.set(sum);
        // Set output with the key
        context.write(key, result);
    }
}
```

1. Mapper implementation and the main method invoker for the question 1 would be,

```
public class Q1 {

    public static class DeliveryMapper
        extends Mapper<Object, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        @Override
        public void map(
            final Object key,
            final Text lineText,
            final Context context
        ) {
            try {
                final String line = lineText.toString();
                final String extraRuns = line.split(",")[8];
                final String isWicket = line.split(",")[11];
                final String totalRuns = line.split(",")[9];

                // If any extra runs were occurred count them
                if (Integer.parseInt(extraRuns) > 0) {
                    word.set("extraRuns");
                    context.write(word, one);
                }

                // If any wicket was taken count them
                if ("1".equals(isWicket)) {
                    word.set("wicketCount");
                    context.write(word, one);
                }

                // If any no runs were observed count them
                if ("0".equals(totalRuns)) {
                    word.set("noRuns");
                    context.write(word, one);
                }
            } catch (Exception e) {
                // There are some issues in the dataset causes number format exception
                // Ignoring them
            }
        }
    }

    public static void main(String[] args) throws Exception {
        // Invoker static method is mostly same as we had in the lab sessions
        // The full implementation can be seen in the zip file
    }
}
```

Now let's call the **yarn** with copied jar file to execute the map reduce job for the question 1 in the environment that we are already in.

## yarn jar coursework-mapreduce-1.0-SNAPSHOT.jar mapreduce.Q1 /ipl /output

```
sh-4.1# yarn jar coursework-mapreduce-1.0-SNAPSHOT.jar mapreduce.Q1 /ipl /output
22/12/05 22:20:35 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
22/12/05 22:20:36 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/12/05 22:20:36 INFO input.FileInputFormat: Total input paths to process : 1
22/12/05 22:20:36 INFO mapreduce.JobSubmitter: number of splits:1
22/12/05 22:20:36 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1670256987498_0011
22/12/05 22:20:36 INFO impl.YarnClientImpl: Submitted application application_1670256987498_0011
22/12/05 22:20:36 INFO mapreduce.Job: The url to track the job: http://hdfs:8088/proxy/application_1670256987498_0011/
22/12/05 22:20:36 INFO mapreduce.Job: Running job: job_1670256987498_0011
22/12/05 22:20:41 INFO mapreduce.Job: Job job_1670256987498_0011 running in uber mode : false
22/12/05 22:20:41 INFO mapreduce.Job: map 0% reduce 0%
22/12/05 22:20:47 INFO mapreduce.Job: map 100% reduce 0%
22/12/05 22:20:53 INFO mapreduce.Job: map 100% reduce 100%
22/12/05 22:20:53 INFO mapreduce.Job: Job job_1670256987498_0011 completed successfully
22/12/05 22:20:53 INFO mapreduce.Job: Counters: 49
File System Counters
  FILE: Number of bytes read=53
  FILE: Number of bytes written=231659
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=22721754
  HDFS: Number of bytes written=46
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=3339
  Total time spent by all reduces in occupied slots (ms)=3077
  Total time spent by all map tasks (ms)=3339
  Total time spent by all reduce tasks (ms)=3077
  Total vcore-seconds taken by all map tasks=3339
  Total vcore-seconds taken by all reduce tasks=3077
  Total megabyte-seconds taken by all map tasks=3419136
  Total megabyte-seconds taken by all reduce tasks=3150848
Map-Reduce Framework
  Map input records=193469
  Map output records=87569
  Map output bytes=1041433
  Map output materialized bytes=53
  Input split bytes=98
  Combine input records=87569
  Combine output records=3
  Reduce input groups=3
  Reduce shuffle bytes=53
  Reduce input records=3
  Reduce output records=3
  Spilled Records=6
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=873
  CPU time spent (ms)=4100
  Physical memory (bytes) snapshot=432775168
  Virtual memory (bytes) snapshot=1728159744
  Total committed heap usage (bytes)=401604608
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
```

Now let's check the output.

## hdfs dfs -cat /output/part-r-00000

```
sh-4.1# ^C
sh-4.1# hdfs dfs -ls /output
22/12/05 22:22:33 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 root supergroup 0 2022-12-05 22:20 /output/_SUCCESS
-rw-r--r-- 1 root supergroup 46 2022-12-05 22:20 /output/part-r-00000
sh-4.1# hdfs dfs -cat /output/part-r-00000
22/12/05 22:22:49 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
extraRuns 10233
noRuns 67841
wicketCount 9495
sh-4.1#
```

Now let's compose the output into a tabular format.

Variable	Description	Value
extraRuns	Extra Runs taken	10233
noRuns	No Runs were taken	67841
wicketCount	Wickets were taken	9495

## 2. Mapper Implementation and the main method for the question 2 would be,

```
public class Q2 {  
  
    public static class WicketsMapper  
        extends Mapper<Object, Text, Text, IntWritable> {  
  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
  
        @Override  
        public void map(  
            final Object key,  
            final Text lineText,  
            final Context context  
        ) {  
            final String line = lineText.toString();  
            try {  
                final String bowlingTeam = line.split(",")[17].toUpperCase().trim();  
                final String isWicket = line.split(",")[11].trim();  
                // If any wicket was taken recorded them with the team  
                if ("1".equals(isWicket.trim())) {  
                    word.set(bowlingTeam);  
                    context.write(word, one);  
                }  
            } catch (Exception e) {  
                e.printStackTrace();  
                // Ignore any exception was occurred  
            }  
        }  
    }  
  
    public static void main(String[] args) throws Exception {  
        // Invoker static method is mostly same as we had in the lab sessions  
        // The full implementation can be seen in the zip file  
    }  
}
```

Now let's execute the Q2 class to obtain results in the jar file.

**yarn jar coursework-mapreduce-1.0-SNAPSHOT.jar mapreduce.Q2 /ipl /output**



```

Deleted /output
sh-4.1# yarn jar coursework-mapreduce-1.0-SNAPSHOT.jar mapreduce.Q2 /ipl /output
22/12/05 22:30:49 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
22/12/05 22:30:49 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/12/05 22:30:50 INFO input.FileInputFormat: Total input paths to process : 1
22/12/05 22:30:50 INFO mapreduce.JobSubmitter: number of splits:1
22/12/05 22:30:50 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1670256987498_0012
22/12/05 22:30:50 INFO impl.YarnClientImpl: Submitted application application_1670256987498_0012
22/12/05 22:30:50 INFO mapreduce.Job: The url to track the job: http://hdfs:8088/proxy/application_1670256987498_0012/
22/12/05 22:30:50 INFO mapreduce.Job: Running job: job_1670256987498_0012
22/12/05 22:30:55 INFO mapreduce.Job: Job job_1670256987498_0012 running in uber mode : false
22/12/05 22:30:55 INFO mapreduce.Job: map 0% reduce 0%
22/12/05 22:31:01 INFO mapreduce.Job: map 100% reduce 0%
22/12/05 22:31:06 INFO mapreduce.Job: map 100% reduce 100%
22/12/05 22:31:06 INFO mapreduce.Job: Job job_1670256987498_0012 completed successfully
22/12/05 22:31:06 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=417
        FILE: Number of bytes written=232387
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=22721754
        HDFS: Number of bytes written=381
        HDFS: Number of read operations=6
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=1
        Launched reduce tasks=1
        Data-local map tasks=1
        Total time spent by all maps in occupied slots (ms)=3075
        Total time spent by all reduces in occupied slots (ms)=2214
        Total time spent by all map tasks (ms)=3075
        Total time spent by all reduce tasks (ms)=2214
        Total vcore-seconds taken by all map tasks=3075
        Total vcore-seconds taken by all reduce tasks=2214
        Total megabyte-seconds taken by all map tasks=3148800
        Total megabyte-seconds taken by all reduce tasks=2267136
    Map-Reduce Framework
        Map input records=193469
        Map output records=9495
        Map output bytes=237089
        Map output materialized bytes=417
        Input split bytes=98
        Combine input records=9495
        Combine output records=16
        Reduce input groups=16
        Reduce shuffle bytes=417
        Reduce input records=16
        Reduce output records=16
        Spilled Records=32
        Shuffled Maps =1
        Failed Shuffles=0
        Merged Map outputs=1
        GC time elapsed (ms)=747
        CPU time spent (ms)=3550
        Physical memory (bytes) snapshot=425197568
        Virtual memory (bytes) snapshot=1727336448
        Total committed heap usage (bytes)=398983168
    Shuffle Errors
        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0

```

And let's see the output in the HDFS.

**hdfs dfs -cat /output/part-r-00000**

```

sh-4.1#
sh-4.1# hdfs dfs -cat /output/part-r-00000
22/12/06 10:07:58 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform.
"CHENNAI SUPER KINGS"      1089
"DECCAN CHARGERS"         454
"DELHI CAPITALS"          212
"DELHI DAREDEVILS"        900
"GUJARAT LIONS"           152
"KINGS XI PUNJAB"         1086
"KOCHI TUSKERS KERALA"    73
"KOLKATA KNIGHT RIDERS"   1068
"MUMBAI INDIANS"          1231
"PUNE WARRIORS"           236
"RAJASTHAN ROYALS"        920
"RISING PUNE SUPERGIANT"   109
"RISING PUNE SUPERGIANTS"  79
"ROYAL CHALLENGERS BANGALORE" 1109
"SUNRISERS HYDERABAD"     759
NA                          18
sh-4.1#

```

And let's tabulate the results were taken,

Team	Wickets
CHENNAI SUPER KINGS	1089
DECCAN CHARGERS	454
DELHI CAPITALS	212
DELHI DAREDEVILS	900
GUJARAT LIONS	152
KINGS XI PUNJAB	1086
KOCHI TUSKERS KERALA	73
KOLKATA KNIGHT RIDERS	1068
MUMBAI INDIANS	1231
PUNE WARRIORS	236
RAJASTHAN ROYALS	920
RISING PUNE SUPERGIANT	109
RISING PUNE SUPERGIANTS	79
ROYAL CHALLENGERS BANGALORE	1109
SUNRISERS HYDERABAD	759
NA	18

## 2.2 Analyze the following using Hive or Pig

For this analysis I am using **Hive** because of its simplicity, and let's create an external table using the dataset that we have put in the **HDFS** using the following command to reuse the **HDFS** storage.

**create external table ipl\_external**

**(id int, inning int, over\_playing int, ball int, batsman string, non\_striker string, bowler string, batsman\_runs int, extra\_runs int, total\_runs int, non\_boundary int, is\_wicket int, dismissal\_kind string, player\_dismissed string, fielder string, extras\_type string, batting\_team string, bowling\_team string) row format delimited fields terminated by ',' LOCATION '/ipl/';**

```
22/12/05 11:59:17 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
sh-4.1# hive

Logging initialized using configuration in jar:file:/usr/local/apache-hive-1.2.2-bin/lib/hive-common-1.2.2.jar!/hive-log4j.properties
hive> create external table ipl_external
  > (id int, inning int, over_playing int, ball int, batsman string, non_striker string, bowler string, batsman_runs int, extra_runs int, total_runs int, non_boundary int, is_wicket int, dismissal_kind string, player_dismissed string, fielder string, extras_type string, batting_team string, bowling_team string)
  > row format delimited fields terminated by ',' LOCATION '/ipl/';
OK
Time taken: 0.652 seconds
hive>
```

1. Let's run the following query to find out the top 10 performing team in the hive terminal.

**select batting\_team, sum(total\_runs) as team\_total\_runs from ipl\_external group by batting\_team order by team\_total\_runs desc limit 10;**

```

Time taken: 0.652 seconds
hive> select batting_team, sum(total_runs) as team_total_runs from ipl_external group by batting_team order by team_total_runs desc limit 10;
Query ID = root_20221205120538_e6b30725-0ba1-42db-aed4-1089ac76dfb8
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1670256987498_0001, Tracking URL = http://hdfs:8088/proxy/application_1670256987498_0001/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1670256987498_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-12-05 12:05:45,620 Stage-1 map = 0%, reduce = 0%
2022-12-05 12:05:50,843 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.62 sec
2022-12-05 12:05:56,130 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.97 sec
MapReduce Total cumulative CPU time: 2 seconds 970 msec
Ended Job = job_1670256987498_0001
Launching Job 2 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1670256987498_0002, Tracking URL = http://hdfs:8088/proxy/application_1670256987498_0002/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1670256987498_0002
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2022-12-05 12:06:06,673 Stage-2 map = 0%, reduce = 0%
2022-12-05 12:06:10,935 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 1.08 sec
2022-12-05 12:06:16,142 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 2.8 sec
MapReduce Total cumulative CPU time: 2 seconds 800 msec
Ended Job = job_1670256987498_0002
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.97 sec HDFS Read: 22730268 HDFS Write: 1130 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 2.8 sec HDFS Read: 5817 HDFS Write: 264 SUCCESS
Total MapReduce CPU Time Spent: 5 seconds 770 msec
OK
"Mumbai Indians" 32259
"Royal Challengers Bangalore" 30183
"Kings XI Punjab" 29990
"Kolkata Knight Riders" 29357
"Chennai Super Kings" 28344
"Rajasthan Royals" 24480
"Delhi Daredevils" 24264
"Sunrisers Hyderabad" 19314
"Deccan Chargers" 11448
"Pune Warriors" 6348
Time taken: 38.333 seconds, Fetched: 10 row(s)
hive>

```

Result gives us the output of top ten teams as follows,

Ranking	Team Name	Score
1	Mumbai Indians	32259
2	Royal Challengers Bangalore	30183
3	Kings XI Punjab	29990
4	Kolkata Knight Riders	29357
5	Chennai Super Kings	28344
6	Rajasthan Royals	24480
7	Delhi Daredevils	24264
8	Sunrisers Hyderabad	19314
9	Deccan Chargers	11448
10	Pune Warriors	6348

- let's run following query in the hive terminal to obtain average runs per over, per inning in the dataset.

**select id, inning, sum(total\_runs)/20 as average\_runs from ipl\_external group by id, inning order by id, inning;**

```
hive> SELECT id, inning, sum(total_runs)/20 as average_runs from ipl_external group by id, inning order by id, inning;
Query ID = root_20221209232758_efedca51-384c-4b4b-be26-7f04b644d1b5
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1670256987498_0034, Tracking URL = http://hdfs:8088/proxy/application_1670256987498_0034/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1670256987498_0034
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-12-09 23:28:03,235 Stage-1 map = 0%, reduce = 0%
2022-12-09 23:28:08,391 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.84 sec
2022-12-09 23:28:14,672 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.26 sec
MapReduce Total cumulative CPU time: 4 seconds 260 msec
Ended Job = job_1670256987498_0034
Launching Job 2 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1670256987498_0035, Tracking URL = http://hdfs:8088/proxy/application_1670256987498_0035/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1670256987498_0035
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2022-12-09 23:28:24,103 Stage-2 map = 0%, reduce = 0%
2022-12-09 23:28:28,239 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 1.56 sec
2022-12-09 23:28:34,447 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 3.41 sec
MapReduce Total cumulative CPU time: 3 seconds 410 msec
Ended Job = job_1670256987498_0035
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.26 sec HDFS Read: 22731603 HDFS Write: 49493 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 3.41 sec HDFS Read: 54149 HDFS Write: 22585 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 670 msec
OK
NULL NULL NULL
335982 1 11.1
335982 2 4.1
335983 1 12.0
335983 2 10.35
335984 1 6.45
335984 2 6.6
335985 1 8.25
335985 2 8.3
335986 1 5.5
335986 2 5.6
335987 1 8.3
335987 2 8.4
335988 1 7.1
335988 2 7.15
335989 1 10.4
335989 2 10.1
335990 1 10.7
335990 2 10.85
335991 1 9.1
335991 2 5.8
335992 1 5.35
```

The result is so lengthy. So, here I have included the topmost part of the output, and the columns are showing as **null** for some reason they should be **id, inning, average score** accordingly but the output result seems accurate.

## 2.3 Analyze the following using Spark.

First of all we need to run the zeppelin detached container with a mounted volume which contains **ipl-data.csv** dataset, and then import the notebook file **CMM-705\_CW\_final.zpln** in the UI of zeppelin web interface which is exposed in local port **8080**

**docker run -it --name zeppelin -p 8080:8080 -v /home/jmadushan/rgu/CMM705/cw/ipl:/data apache/zeppelin:0.9.0**

```
madushan@CLLK-JANITHAM:~$ docker run -it --name zeppelin -d -p 8080:8080 -v /home/jmadushan/rgu/CMM705/cw/ipl:/data apache/zeppelin:0.9.0
d14996dad36422c2be8fd3bac5df22cc5768d897c186725e54c96ccddc5d6f1
madushan@CLLK-JANITHAM:~$
```

First, we need to import the dataset from the mounted location and load it into the data frame.

```
df = spark.read.option("header",True).csv("/data/ipl-data.csv")
```

```
df.show(10)
```

```
df.createOrReplaceTempView("batting")
```

```
%spark
# Load CSV file into the dataframe
df = spark.read.option("header",True).csv("/data/ipl-data.csv")

df.show(10)

# Create a view called batting
df.createOrReplaceTempView("batting")
```

	id	inning	lover	ball	batsman	non_striker	bowler	batsman_runs	extra_runs	total_runs	non_boundary	is_wicket	dismissal_kind	player_dismissed	fielder	extras_type	batting_team	bowling_team
[335982]	1	6	5	RT Ponting	BB McCullum	AA NooFke		1	0	1	0	0	0	NA	NA	NA	Kolkata Knight Ri...	Royal Challengers...
[335982]	1	6	6	BB McCullum	RT Ponting	AA NooFke		1	0	1	0	0	0	NA	NA	NA	Kolkata Knight Ri...	Royal Challengers...
[335982]	1	7	1	BB McCullum	RT Ponting	Z Khan		0	0	0	0	0	0	NA	NA	NA	Kolkata Knight Ri...	Royal Challengers...
[335982]	1	7	2	BB McCullum	RT Ponting	Z Khan		1	0	1	0	0	0	NA	NA	NA	Kolkata Knight Ri...	Royal Challengers...
[335982]	1	7	3	RT Ponting	BB McCullum	Z Khan		1	0	1	0	0	0	NA	NA	NA	Kolkata Knight Ri...	Royal Challengers...
[335982]	1	7	4	BB McCullum	RT Ponting	Z Khan		1	0	1	0	0	0	NA	NA	NA	Kolkata Knight Ri...	Royal Challengers...
[335982]	1	7	5	RT Ponting	BB McCullum	Z Khan		1	0	1	0	0	0	NA	NA	NA	Kolkata Knight Ri...	Royal Challengers...
[335982]	1	7	6	BB McCullum	RT Ponting	Z Khan		1	0	1	0	0	0	NA	NA	NA	Kolkata Knight Ri...	Royal Challengers...
[335982]	1	8	1	BB McCullum	RT Ponting	JH Kallis		0	0	0	0	0	0	NA	NA	NA	Kolkata Knight Ri...	Royal Challengers...
[335982]	1	8	2	BB McCullum	RT Ponting	JH Kallis		0	0	0	0	0	0	NA	NA	NA	Kolkata Knight Ri...	Royal Challengers...

only showing top 10 rows

1. In the above step we have created a view called **`batting`** on the main dataset, and now let's filter out/aggregate the data to give an output of **each player's total score per inning**.

**# Create a summary query to find total runs of each team member for an inning**

```
inning_df = spark.sql("SELECT id, inning, batsman, sum(total_runs) as total_runs_per_inning FROM batting group by id, inning, batsman")
```

```
inning_df.show(10)
```

```
%spark

# Create a summary query to find total runs of each team member for an inning
inning_df = spark.sql("SELECT id, inning, batsman, sum(total_runs) as total_runs_per_inning FROM batting group by id, inning, batsman")
inning_df.show(10)
```

	id	inning	batsman	total_runs_per_inning
[336029]	1	P Kumar	21.0	
[336037]	1	SB Styris	24.0	
[392190]	2	CH Gayle	48.0	
[392196]	1	JDP Oram	42.0	
[392208]	2	RV Uthappa	68.0	
[392218]	2	R Dravid	12.0	
[392220]	1	DR Smith	48.0	
[392230]	2	MM Patel	23.0	
[419124]	2	MS Bisla	37.0	
[419128]	2	R Dravid	1.0	

only showing top 10 rows

Took 1 sec. Last updated by anonymous at December 10 2022, 1:43:59 PM.

Now let's summarize **the first question** by performing operations on RDD of above created summarized `inning_df` data frame.

**# Count total number of players**

```
total_players = df.rdd.map(lambda row: row.batsman).distinct().count()
```

**# Count total number of players scored more than 50 per inning**

```
players_have_scored_more_than_50 = inning_df.rdd.filter(lambda p:  
p.total_runs_per_inning>50).map(lambda inning: inning.batsman).distinct().count()
```

```
print("Total number of players: %d\n" % total_players)
```

```
print("Number of players who have scored more than 50 per inning is : %d\n" %  
players_have_scored_more_than_50)
```

```
print("Percentage of players who have scored more than 50 per inning is : %.2f " %  
(players_have_scored_more_than_50/total_players*100))
```

```
%pyspark  
  
# Count total number of players  
total_players = df.rdd.map(lambda row: row.batsman).distinct().count()  
  
# Count total number of players scored more than 50 per inning  
players_have_scored_more_than_50 = inning_df.rdd.filter(lambda p: p.total_runs_per_inning>50).map(lambda inning: inning.batsman).distinct().count()  
  
print("Total number of players: %d\n" % total_players)  
print("Number of players who have scored more than 50 per inning is : %d\n" % players_have_scored_more_than_50)  
print("Percentage of players who have scored more than 50 per inning is : %.2f " % (players_have_scored_more_than_50/total_players*100))  
  
Total number of players: 537  
  
Number of players who have scored more than 50 per inning is : 157  
  
Percentage of players who have scored more than 50 per inning is : 29.24  
  
Took 4 sec. Last updated by anonymous at December 10 2022, 1:44:11 PM.
```

Results obtained by the above analysis,

Description	Result
Total Number of distinct Players	537
Total Number of Players have scored more than 50 in a single inning	157
Percentage of players who have scored more than 50	29.24

- Let's filter out the necessary fields from the main dataset in the data frame **df** and filter out the only the fields that are required to answer this question (**id**, **total\_runs**, **batting\_team**, **bowling\_team**).

```
selected_df = df.select(df['id'], df['total_runs'], df['batting_team'], df['bowling_team'])
```

```
selected_df.show(10)
```

```
%pyspark
# First select only the fields that's need for the analysis from the original dataframe
selected_df = df.select(df['id'], df['total_runs'], df['batting_team'], df['bowling_team'])
selected_df.show(10)
```

	id	total_runs	batting_team	bowling_team
	335982	1	Kolkata Knight Ri...	Royal Challengers...
	335982	1	Kolkata Knight Ri...	Royal Challengers...
	335982	0	Kolkata Knight Ri...	Royal Challengers...
	335982	1	Kolkata Knight Ri...	Royal Challengers...
	335982	1	Kolkata Knight Ri...	Royal Challengers...
	335982	1	Kolkata Knight Ri...	Royal Challengers...
	335982	1	Kolkata Knight Ri...	Royal Challengers...
	335982	0	Kolkata Knight Ri...	Royal Challengers...
	335982	0	Kolkata Knight Ri...	Royal Challengers...

only showing top 10 rows

Now let's create a view called '**match**' on the data frame **selected\_df**, and let's perform an SQL operation to find out total runs of each team as batting team in the each match.

```
# create batting view called `match` for the selected columns
```

```
selected_df.createOrReplaceTempView("match")
```

```
# create summary dataframe from each match and scores for batting and bowling teams.
```

```
match_summary_df = spark.sql("SELECT id, sum(total_runs) as total_runs, batting_team, bowling_team FROM match group by id, batting_team, bowling_team order by id")
```

```
match_summary_df.show(10)
```

```
%pyspark
# create batting view called `match` for the selected columns
selected_df.createOrReplaceTempView("match")

# create summary dataframe from each match and scores for batting and bowling teams.
match_summary_df = spark.sql("SELECT id, sum(total_runs) as total_runs, batting_team, bowling_team FROM match group by id, batting_team, bowling_team order by id")
match_summary_df.show(10)
```

	id	total_runs	batting_team	bowling_team
	1082591	172.0	Royal Challengers...	Sunrisers Hyderabad
	1082591	207.0	Sunrisers Hyderabad	Royal Challengers...
	1082592	184.0	Mumbai Indians	Rising Pune Super...
	1082592	187.0	Rising Pune Super...	Mumbai Indians
	1082593	184.0	Kolkata Knight Ri...	Gujarat Lions
	1082593	183.0	Gujarat Lions	Kolkata Knight Ri...
	1082594	164.0	Kings XI Punjab	Rising Pune Super...
	1082594	163.0	Rising Pune Super...	Kings XI Punjab
	1082595	157.0	Royal Challengers...	Delhi Daredevils
	1082595	142.0	Delhi Daredevils	Royal Challengers...

only showing top 10 rows

In the above output we can see match id is duplicated by 2 rows, means for each match it contains rows for each team as **batting team** and corresponding **bowling team**.

Row => [id], [total\_runs\_team\_A, total\_runs\_team\_B], [Team\_A, Team\_B]

And now let's perform aggregate function to aggregate columns to describe each match by a single row. So, aggregate functions have been called **total\_runs**, **batting\_team**, and **bowling\_team** column as it is explained above.

from pyspark.sql.functions import collect\_list, col

# Now lets aggregate scores and corresponding teams of each match, and this would be the base dataframe to answer the questions.

```
match_scores_df =
match_summary_df.groupBy("id").agg(collect_list(col("total_runs")).alias("scores"),
collect_list(col("batting_team")).alias("teams"))
```

match\_scores\_df.show(10)

```
%pyspark
from pyspark.sql.functions import collect_list, col

# Now lets aggregate scores and corresponding teams of each match, and this would be the base dataframe to answer the questions.
match_scores_df = match_summary_df.groupBy("id").agg(collect_list(col("total_runs")).alias("scores"), collect_list(col("batting_team")).alias("teams"))
match_scores_df.show(10)
```

id	scores	teams
1082591	[207.0, 172.0]	[Sunrisers Hydera...
1082592	[187.0, 184.0]	[Rising Pune Supe...
1082593	[184.0, 183.0]	[Kolkata Knight R...
1082594	[164.0, 163.0]	[Kings XI Punjab,...]
1082595	[142.0, 157.0]	[Delhi Daredevils...
1082596	[140.0, 135.0]	[Sunrisers Hydera...
1082597	[180.0, 178.0]	[Mumbai Indians, ...]
1082598	[150.0, 148.0]	[Kings XI Punjab,...]
1082599	[205.0, 108.0]	[Delhi Daredevils...
1082600	[158.0, 159.0]	[Sunrisers Hydera...

only showing top 10 rows

Analyze data for won matches.

And now let's add an additional column mentioning **which team has won** each game.

from pyspark.sql import functions as f

```
winners_df=match_scores_df[match_scores_df['scores'][0]!=match_scores_df['scores'][1]].withColumn('winners', f.when(f.col('scores')[0] > f.col('scores')[1],
f.col('teams')[0]).otherwise(f.col('teams')[1]))
```

winners\_df.show(10)

```
%pyspark
from pyspark.sql import functions as f

# Now lets compare the batting scores and add winners column
winners_df=match_scores_df[match_scores_df['scores'][0]!=match_scores_df['scores'][1]].withColumn('winners', f.when(f.col('scores')[0] > f.col('scores')[1], f.col('teams')[0]).otherwise(f.col('teams')[1]))
winners_df.show(10)
```

id	scores	teams	winners
1082591	[207.0, 172.0]	[Sunrisers Hydera...	Sunrisers Hyderabad
1082592	[187.0, 184.0]	[Rising Pune Supe...	Rising Pune Super...
1082593	[184.0, 183.0]	[Kolkata Knight R...	Kolkata Knight Ri...
1082594	[164.0, 163.0]	[Kings XI Punjab,...]	Kings XI Punjab
1082595	[142.0, 157.0]	[Delhi Daredevils...	Royal Challengers...
1082596	[140.0, 135.0]	[Sunrisers Hydera...	Sunrisers Hyderabad
1082597	[180.0, 178.0]	[Mumbai Indians, ...]	Mumbai Indians
1082598	[150.0, 148.0]	[Kings XI Punjab,...]	Kings XI Punjab
1082599	[205.0, 108.0]	[Delhi Daredevils...	Delhi Daredevils
1082600	[158.0, 159.0]	[Sunrisers Hydera...	Mumbai Indians

only showing top 10 rows



Finally, let's create a view called **winners** on the **winners\_df** that we created in the previous step, and summarize each team's total number of winning matches.

```
winners_df.createOrReplaceTempView("winners")
```

```
winners_summary_df = spark.sql("SELECT winners as team, count(*) as wins FROM winners group by winners")
```

```
winners_summary_df.show(10)
```

```
%pyspark
# Create view as winners
winners_df.createOrReplaceTempView("winners")

# Summarize each team's winning count
winners_summary_df = spark.sql("SELECT winners as team, count(*) as wins FROM winners group by winners")
winners_summary_df.show(10)
```

team	wins
Sunrisers Hyderabad	67
Chennai Super Kings	106
Rising Pune Super...	10
Deccan Chargers	29
Kochi Tuskers Kerala	6
Rajasthan Royals	79
Gujarat Lions	13
Royal Challengers...	89
Kolkata Knight Ri...	96
Rising Pune Super...	4

only showing top 10 rows

### Analyze data for matches lost.

Next, let's create another data frame called **losers\_df** containing an additional column mentioning the losing team.

```
losers_df=match_scores_df[match_scores_df['scores'][0]!=match_scores_df['scores'][1]].withColumn('losers', f.when(f.col('scores')[0] > f.col('scores')[1], f.col('teams')[1]).otherwise(f.col('teams')[0]))
```

```
losers_df.show(10)
```

```
%pyspark
# Summarize loosing team for each match
losers_df=match_scores_df[match_scores_df['scores'][0]!=match_scores_df['scores'][1]].withColumn('losers', f.when(f.col('scores')[0] > f.col('scores')[1], f.col('teams')[1]).otherwise(f.col('teams')[0]))
losers_df.show(10)
```

id	scores	teams	losers
1082591	[207.0, 172.0]	[Sunrisers Hydera..., Royal Challengers...]	Royal Challengers...
1082592	[187.0, 184.0]	[Rising Pune Supe..., Mumbai Indians]	Mumbai Indians
1082593	[184.0, 183.0]	[Kolkata Knight R..., Gujarat Lions]	Gujarat Lions
1082594	[164.0, 163.0]	[Kings XI Punjab, ..., Rising Pune Super...]	Rising Pune Super...
1082595	[142.0, 157.0]	[Delhi Daredevils..., Delhi Daredevils]	Delhi Daredevils
1082596	[140.0, 135.0]	[Sunrisers Hydera..., Gujarat Lions]	Gujarat Lions
1082597	[180.0, 178.0]	[Mumbai Indians, ..., Kolkata Knight Ri...]	Kolkata Knight Ri...
1082598	[150.0, 148.0]	[Kings XI Punjab, ..., Royal Challengers...]	Royal Challengers...
1082599	[205.0, 108.0]	[Delhi Daredevils..., Rising Pune Super...]	Rising Pune Super...
1082600	[158.0, 159.0]	[Sunrisers Hydera..., Sunrisers Hyderabad]	Sunrisers Hyderabad

only showing top 10 rows

And create a view called losers mentioning each team losing matches.

```
losers_df.createOrReplaceTempView("losers")
```

```
losers_summary_df = spark.sql("SELECT losers as team, count(*) as losses FROM losers group by losers")
```

**loosers\_summary\_df.show(10)**

```
%pyspark
# create a view called losers
loosers_df.createOrReplaceTempView("losers")

# summarize losing count for each team
loosers_summary_df = spark.sql("SELECT losers as team, count(*) as losses FROM losers group by losers")
loosers_summary_df.show(10)
```

team	losses
Sunrisers Hyderabad	54
Chennai Super Kings	71
Rising Pune Super...	6
Deccan Chargers	46
Kochi Tuskers Kerala	8
Rajasthan Royals	78
Gujarat Lions	16
Royal Challengers...	102
Kolkata Knight Ri...	92
Rising Pune Super...	10

only showing top 10 rows

**Analyze data for matches drawn.**

And now let's find the occurrences that each team has drawn.

**eq\_df=match\_scores\_df[match\_scores\_df['scores'][0]==match\_scores\_df['scores'][1]]**

```
%pyspark

# filter the matches has been drawn
eq_df=match_scores_df[match_scores_df['scores'][0]==match_scores_df['scores'][1]]

eq_selected_df = eq_df.select(['teams'])

# Splitup the teams of lists
eq_selected_rdd=eq_selected_df.rdd.flatMap(lambda x:x).flatMap(lambda x:x)

# create a dataframe as draw with splitted team names
draw_df = eq_selected_rdd.map(lambda x: (x, )).toDF(["draw"])
draw_df.show(10)
```

draw
Gujarat Lions
Mumbai Indians
Delhi Capitals
Kolkata Knight Ri...
Sunrisers Hyderabad
Mumbai Indians
Kings XI Punjab
Delhi Capitals
Kolkata Knight Ri...
Sunrisers Hyderabad

only showing top 10 rows

**eq\_selected\_df = eq\_df.select(['teams'])**

**eq\_selected\_rdd=eq\_selected\_df.rdd.flatMap(lambda x:x).flatMap(lambda x:x)**

**draw\_df = eq\_selected\_rdd.map(lambda x: (x, )).toDF(["draw"])**

**draw\_df.show(10)**

And now let's create a view on **drawn\_df** called **`draw`** to summarize the total number of matches that have been drawn by each team.

**draw\_df.createOrReplaceTempView("draws")**

```
draw_summary_df = spark.sql("SELECT draw as team, count(*) as draws FROM draws group by draw")
draw_summary_df.show(10)
```

```
%pyspark
# Create a dataframe called draws
draw_df.createOrReplaceTempView("draws")

# Summarize each teams drawn count
draw_summary_df = spark.sql("SELECT draw as team, count(*) as draws FROM draws group by draw")
draw_summary_df.show(10)
```

team	draws
Sunrisers Hyderabad	3
Chennai Super Kings	1
Rajasthan Royals	3
Gujarat Lions	1
Royal Challengers...	3
Kolkata Knight Ri...	4
Kings XI Punjab	4
Delhi Daredevils	1
Delhi Capitals	2
Mumbai Indians	4

And let's aggregate all outputs into a single data frame to summarize each team's performance.

```
summary_df=winners_summary_df.join(loosers_summary_df, ["team"]).join(draw_summary_df, ["team"])
```

```
summary_df.show()
```

```
%pyspark
# Aggregate all the dataframes together to summarize teach teams statistics
summary_df=winners_summary_df.join(loosers_summary_df, ["team"]).join(draw_summary_df, ["team"])
summary_df.show()
```

team	wins	losses	draws
Sunrisers Hyderabad	67	54	3
Chennai Super Kings	106	71	1
Rajasthan Royals	79	78	3
Gujarat Lions	13	16	1
Royal Challengers...	89	102	3
Kolkata Knight Ri...	96	92	4
Kings XI Punjab	85	101	4
Delhi Daredevils	70	89	1
Delhi Capitals	17	14	2
Mumbai Indians	118	81	4

Finally, let's export **summary\_df** into a csv file to import in the dashboard.

```
summary_df.coalesce(1).write.options(header='True', delimiter=',')
.format("csv").mode('overwrite').csv("/data/summary")
```

### 3. Performing Machine Learning model using Spark MLlib

We have already loaded the dataset into a data frame **df** and created a view called **batting** and let's fetch the first ten columns of that view.

```
spark.sql("SELECT * FROM batting").show(10)
```

```
%pyspark
spark.sql("SELECT * FROM batting").show(10)
```

	id	inning	over	ball	batsman	non_striker	bowler	batsman_runs	extra_runs	total_runs	non_boundary	is_wicket	dismissal_kind	player_dismissed	fielder	extras_type	betting_team	bowling_team
[335982]	1	6	5	RT Ponting	BB McCullum	AA Noffke		1	0	1	0	0		NA	NA	NA	NA	Kolkata Knight Ri... Royal Challengers...
[335982]	1	6	6	BB McCullum	RT Ponting	AA Noffke		1	0	1	0	0		NA	NA	NA	NA	Kolkata Knight Ri... Royal Challengers...
[335982]	1	7	1	BB McCullum	RT Ponting	Z Khan		0	0	0	0	0		NA	NA	NA	NA	Kolkata Knight Ri... Royal Challengers...
[335982]	1	7	2	BB McCullum	RT Ponting	Z Khan		1	0	1	0	0		NA	NA	NA	NA	Kolkata Knight Ri... Royal Challengers...
[335982]	1	7	3	RT Ponting	BB McCullum	Z Khan		1	0	1	0	0		NA	NA	NA	NA	Kolkata Knight Ri... Royal Challengers...
[335982]	1	7	4	BB McCullum	RT Ponting	Z Khan		1	0	1	0	0		NA	NA	NA	NA	Kolkata Knight Ri... Royal Challengers...
[335982]	1	7	5	RT Ponting	BB McCullum	Z Khan		1	0	1	0	0		NA	NA	NA	NA	Kolkata Knight Ri... Royal Challengers...
[335982]	1	7	6	BB McCullum	RT Ponting	Z Khan		1	0	1	0	0		NA	NA	NA	NA	Kolkata Knight Ri... Royal Challengers...
[335982]	1	8	1	BB McCullum	RT Ponting	JH Kallis		0	0	0	0	0		NA	NA	NA	NA	Kolkata Knight Ri... Royal Challengers...
[335982]	1	8	2	BB McCullum	RT Ponting	JH Kallis		0	0	0	0	0		NA	NA	NA	NA	Kolkata Knight Ri... Royal Challengers...

only showing top 10 rows

let's analyze **overs** column,

```
spark.sql("SELECT distinct over FROM batting order by over asc").show()
```

```
%pyspark
spark.sql("SELECT distinct over FROM batting order by over asc").show()
```

```
+-----+
|over|
+-----+
| 0|
| 1|
| 10|
| 11|
| 12|
| 13|
| 14|
| 15|
| 16|
| 17|
| 18|
| 19|
| 2|
| 3|
| 4|
```

Here we can see the overs are starting from zero. So, I assume **0** as the **first over**, and sum up each team's total score over the matches as batting teams in the first 6 overs. The assignment says Build a model that predicts the average runs expected to be scored in the first 6 overs if **TeamA** team plays against **TeamB**. So, this means we need to select **batting\_team**, **bowling\_team** as features, and sum of runs as the target variable.

```
batting_summary_df = spark.sql("SELECT id, batting_team, bowling_team, sum(total_runs) as sum
FROM batting where over < 6 group by id, batting_team, bowling_team order by id")
```

```
batting_summary_df.show(10)
```

```
%pyspark
batting_summary_df = spark.sql("SELECT id, batting_team, bowling_team, sum(total_runs) as sum FROM batting where over < 6 group by id, batting_team, bowling_team order by id")
batting_summary_df.show(10)
```

id	batting_team	bowling_team	sum
1082591	Sunrisers Hyderabad	Royal Challengers...	59.0
1082591	Royal Challengers...	Sunrisers Hyderabad	54.0
1082592	Rising Pune Super...	Mumbai Indians	59.0
1082592	Mumbai Indians	Rising Pune Super...	61.0
1082593	Gujarat Lions	Kolkata Knight Ri...	52.0
1082593	Kolkata Knight Ri...	Gujarat Lions	73.0
1082594	Kings XI Punjab	Rising Pune Super...	56.0
1082594	Rising Pune Super...	Kings XI Punjab	35.0
1082595	Delhi Daredevils	Royal Challengers...	43.0
1082595	Royal Challengers...	Delhi Daredevils	41.0

only showing top 10 rows

Took 1 sec. Last updated by anonymous at December 10 2022, 4:12:02 PM.

Ok, now we have obtained each team score for the first 6 overs along with the bowling team, but in machine learning, it is necessary to code categorical variables with numeric interpretation, called categorical coding. So, let's perform categorical coding over **batting\_team**, and bowling team variables as **batting\_coded** & **bowling\_coded** variables.

```
from pyspark.ml.feature import StringIndexer, OneHotEncoder

# create indexers

batting_indexer=StringIndexer(inputCol='batting_team', outputCol='batting_coded')

bowling_indexer=StringIndexer(inputCol='bowling_team', outputCol='bowling_coded')

#perform categorical coding

indexer_batting = batting_indexer.fit(batting_summary_df)

indexed_df = indexer_batting.transform(batting_summary_df)

indexer_bowling = bowling_indexer.fit(indexed_df)

indexed_df = indexer_bowling.transform(indexed_df)

indexed_df.show()
```

```
%pyspark
#https://masum-math8065.medium.com/how-to-convert-categorical-data-into-numeric-in-pyspark-2202407f5fac
from pyspark.ml.feature import StringIndexer, OneHotEncoder

# create indexers
batting_indexer=StringIndexer(inputCol='batting_team', outputCol='batting_coded')
bowling_indexer=StringIndexer(inputCol='bowling_team', outputCol='bowling_coded')

#perform categorical coding
indexer_batting = batting_indexer.fit(batting_summary_df)
indexed_df = indexer_batting.transform(batting_summary_df)
indexer_bowling = bowling_indexer.fit(indexed_df)
indexed_df = indexer_bowling.transform(indexed_df)

indexed_df.show()
```

id	batting_team	bowling_team	sum	batting_coded	bowling_coded
392211	Kolkata Knight Ri...	Delhi Daredevils	54.0	2.0	6.0
419148	Royal Challengers...	Kolkata Knight Ri...	48.0	1.0	2.0
501218	Pune Warriors	Mumbai Indians	33.0	9.0	0.0
1082643	Gujarat Lions	Sunrisers Hyderabad	61.0	11.0	7.0
1178429	Sunrisers Hyderabad	Royal Challengers...	52.0	7.0	1.0
1216531	Kings XI Punjab	Royal Challengers...	56.0	3.0	1.0
392212	Deccan Chargers	Mumbai Indians	37.0	8.0	0.0
392238	Royal Challengers...	Chennai Super Kings	59.0	1.0	4.0
501202	Kings XI Punjab	Pune Warriors	36.0	3.0	9.0
501219	Kochi Tuskers Kerala	Kolkata Knight Ri...	45.0	13.0	2.0
501234	Kings XI Punjab	Kolkata Knight Ri...	40.0	3.0	2.0
598002	Chennai Super Kings	Mumbai Indians	36.0	4.0	0.0
598036	Royal Challengers...	Rajasthan Royals	55.0	1.0	5.0
598045	Royal Challengers...	Kings XI Punjab	31.0	1.0	3.0

Now let's filter out only the necessary variables in the **indexed\_df**.

```
coded_df=indexed_df.select("batting_coded", "bowling_coded", "sum")
```

```
coded_df.show()
```

```
%pyspark
coded_df=indexed_df.select("batting_coded", "bowling_coded", "sum")
coded_df.show()
```

batting_coded	bowling_coded	sum
2.0	6.0	54.0
1.0	2.0	48.0
9.0	0.0	33.0
11.0	7.0	61.0
7.0	1.0	52.0
3.0	1.0	56.0
8.0	0.0	37.0
1.0	4.0	59.0
3.0	9.0	36.0
13.0	2.0	45.0
3.0	2.0	40.0
4.0	0.0	36.0
1.0	5.0	55.0
1.0	3.0	31.0

Now let's prepare our final dataset to perform machine learning operations. I am using **VectorAssembler** in **PySpark** to create feature vectors.

```
from pyspark.ml.regression import LinearRegression
```

```
from pyspark.ml.feature import VectorAssembler
```

```
assembler = VectorAssembler(inputCols=[  
    'batting_coded',  
    'bowling_coded'],
```

```
    outputCol='features')
```

```
data_set = assembler.transform(coded_df)
```

```
data_set.show(2)
```

```
%pyspark  
from pyspark.ml.regression import LinearRegression  
from pyspark.ml.feature import VectorAssembler  
  
assembler = VectorAssembler(inputCols=[  
    'batting_coded',  
    'bowling_coded'],  
    outputCol='features')  
  
data_set = assembler.transform(coded_df)  
data_set.show(2)  
  
+-----+-----+-----+-----+  
|batting_coded|bowling_coded|sum|features|  
+-----+-----+-----+-----+  
|          2.0|          6.0|54.0|[2.0,6.0]|  
|          1.0|          2.0|48.0|[1.0,2.0]|  
+-----+-----+-----+-----+  
only showing top 2 rows
```

Let's filter out only the necessary fields from the above dataset.

```
%pyspark  
  
final_df=data_set.select(['features','sum'])  
final_df.show()  
  
+-----+-----+  
|features|sum|  
+-----+-----+  
|[2.0,6.0]|54.0|  
|[1.0,2.0]|48.0|  
|[9.0,0.0]|33.0|  
|[11.0,7.0]|61.0|  
|[7.0,1.0]|52.0|  
|[3.0,1.0]|56.0|  
|[8.0,0.0]|37.0|  
|[1.0,4.0]|59.0|  
|[3.0,9.0]|36.0|  
|[13.0,2.0]|45.0|  
|[3.0,2.0]|40.0|  
|[4.0,0.0]|36.0|  
|[1.0,5.0]|55.0|  
|[1.0,3.0]|31.0|  
|  0.0,0.0]|27.0|
```

And now we need to split our final dataset into 80% as training dataset and 20% as the testing dataset.

```
train_data,test_data = final_df.randomSplit([0.8,0.2])
```

```
%pyspark
train_data,test_data = final_df.randomSplit([0.8,0.2])

Took 0 sec. Last updated by anonymous at December 05 2022, 9:15:41 PM.
```

I am going to use **Linear Regression model** in this case and fit in the `train_data` dataset to train the model.

```
from pyspark.ml.regression import LinearRegression
```

```
linearRegression = LinearRegression(labelCol='sum')
```

```
linearModel = linearRegression.fit(train_data)
```

```
%pyspark
from pyspark.ml.regression import LinearRegression

linearRegression = LinearRegression(labelCol='sum')
linearModel = linearRegression.fit(train_data)

Took 3 sec. Last updated by anonymous at December 05 2022, 9:19:16 PM.
```

Let us obtain parameters to evaluate the model accuracy, such as RMSE and R2 with the `test_data` dataset.

```
test_stats = linearModel.evaluate(test_data)
```

```
print("RMSE:", test_stats.rootMeanSquaredError)
```

```
print("R2: ", test_stats.r2)
```

```
%pyspark
test_stats = linearModel.evaluate(test_data)
print("RMSE:", test_stats.rootMeanSquaredError)
print("R2: ", test_stats.r2)

RMSE: 11.84618875859852
R2: 0.0016427383226924608

Took 3 sec. Last updated by anonymous at December 10 2022, 4:10:24 PM.
```

Until now we have trained the model, but now we must make the prediction for the given scenario. We need to predict what would be the expected score in the first 6 overs if **Mumbai Indians** bat against **Kolkata Night Riders**. So, we need to find corresponding feature matrix.



**indexed\_df.show(5)**

```
%pyspark
indexed_df.show(5)
```

	id	batting_team	bowling_team	sum	batting_coded	bowling_coded
	392211	Kolkata Knight Ri...	Delhi Daredevils	54.0	2.0	6.0
	419148	Royal Challengers...	Kolkata Knight Ri...	48.0	1.0	2.0
	501218	Pune Warriors	Mumbai Indians	33.0	9.0	0.0
	1082643	Gujarat Lions	Sunrisers Hyderabad	61.0	11.0	7.0
	1178429	Sunrisers Hyderabad	Royal Challengers...	52.0	7.0	1.0

only showing top 5 rows

Took 1 sec. Last updated by anonymous at December 05 2022, 9:23:00 PM.

**data\_set.show(5)**

```
%pyspark
data_set.show(5)
```

	batting_coded	bowling_coded	sum	features
	2.0	6.0	54.0	[2.0,6.0]
	1.0	2.0	48.0	[1.0,2.0]
	9.0	0.0	33.0	[9.0,0.0]
	11.0	7.0	61.0	[11.0,7.0]
	7.0	1.0	52.0	[7.0,1.0]

only showing top 5 rows

Took 0 sec. Last updated by anonymous at December 05 2022, 9:23:03 PM.

Now we need to filter out corresponding feature vector for `Mumbai Indians` and `Kolkata Night Riders` Therefore, we inner join `indexed\_df`, and `data\_set` data frames to get the corresponding feature vector.

```
cond = [indexed_df.batting_coded == data_set.batting_coded, indexed_df.bowling_coded == data_set.bowling_coded]
```

```
joined_df=data_set.join(indexed_df, cond, 'inner')
```

```
predict_df=joined_df[joined_df.batting_team=='Mumbai Indians'][joined_df.bowling_team=='Kolkata Knight Riders'].select(['features'])
```

**predict\_df.show(1)**

```
%pyspark
# Now we need to filter out corresponding feature vector for `Mumbai Indians` and `Kolkata Night Riders`
# Therefore, we inner join `indexed_df`, and `data_set` dataframes to get the corresponding feature vector
cond = [indexed_df.batting_coded == data_set.batting_coded, indexed_df.bowling_coded == data_set.bowling_coded]
joined_df=data_set.join(indexed_df, cond, 'inner')
predict_df=joined_df[joined_df.batting_team=='Mumbai Indians'][joined_df.bowling_team=='Kolkata Knight Riders'].select(['features'])
predict_df.show(1)

+-----+
| features|
+-----+
|[0.0,2.0]|
+-----+
only showing top 1 row
```

Took 2 sec. Last updated by anonymous at December 05 2022, 9:11:07 PM.

Predict using the Linear regression model that we have already trained.

**prediction = linearModel.transform(predict\_df)**

**prediction.show(1)**

```
%pyspark
prediction = linearModel.transform(predict_df)
prediction.show(1)

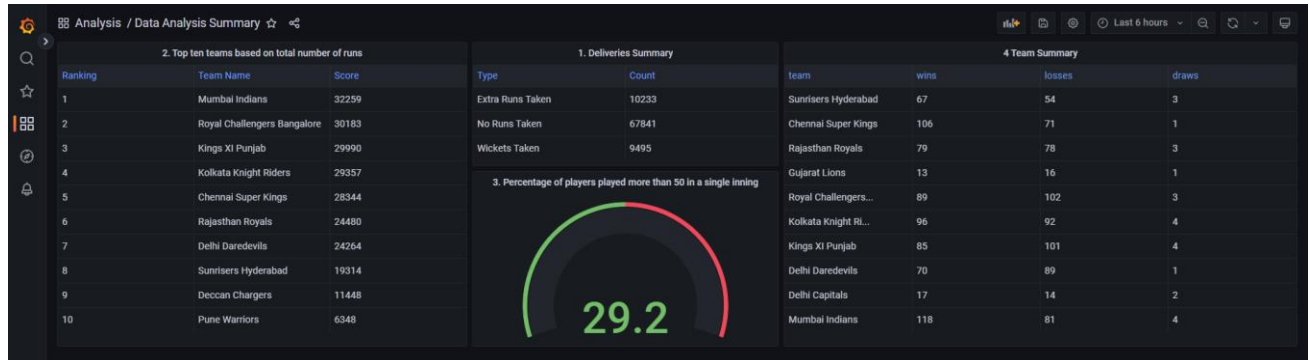
+-----+-----+
| features| prediction|
+-----+-----+
|[0.0,2.0]|45.44378008259662|
+-----+-----+
only showing top 1 row
```

Took 2 sec. Last updated by anonymous at December 05 2022, 9:11:09 PM.

Here we can see the predicted score for **Mumbai Indians** as the batting team, and **Kolkata Knight Riders** as the bowling team is **45.44**

## 4. Presentation of analysis

For the presentation for the analysis, I used Grafana and Infinity plugin to create a static datastore for the dashboard.



### How to run this Dashboard?

In the zip file there is a file called **analysis-dashboard.yaml**, to run this file we need to have Kubernetes cli which is **kubectl** and working cluster. When they are ready, we can simply create the necessary resources using the following command.

**kubectl create -f analysis-dashboard.yaml**

```
jmadushan@CLLK-JANITHAM:~/rgu/CMM705/cw$ kubectl create -f analysis-dashboard.yaml
configmap/analysis-dashboard created
configmap/grafana-dashboards created
configmap/grafana-infinity created
pod/grafana created
jmadushan@CLLK-JANITHAM:~/rgu/CMM705/cw$
```

Let me describe briefly about **analysis-dashboard.yaml** file, it contains following **Kubernetes** resources with corresponding usages.

Resource	Usage
ConfigMap (analysis-dashboard)	Contains dashboard json file which is the json representation of our dashboard
ConfigMap (Grafana-dashboards)	Dashboard provisioners to automate provisioning
ConfigMap (Grafana-infinity)	Data source provisioners to automate infinity data source provisioning
Pod (Grafana)	Grafana dashboard pod containing the Grafana container

And this is a fully automated static dashboard which will automatically provisioning the resources that is necessary. **Configuration as code (CAC)** paradigm has been used here.

And to access this dashboard we need to create a bonded port with local host. For that we could use the following command to open a port locally.

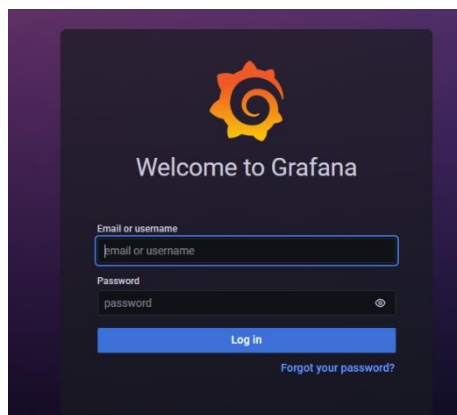
**kubectrl port-forward po/grafana 3000:3000**

```
Forwarding from [::1]:3000 -> 3000
^Cjmadushan@CLLK-JANITHAM:~/rgu/CMM705/cw$ kubectrl port-forward po/grafana 3000:3000
Forwarding from 127.0.0.1:3000 -> 3000
Forwarding from [::1]:3000 -> 3000
Handling connection for 3000
```

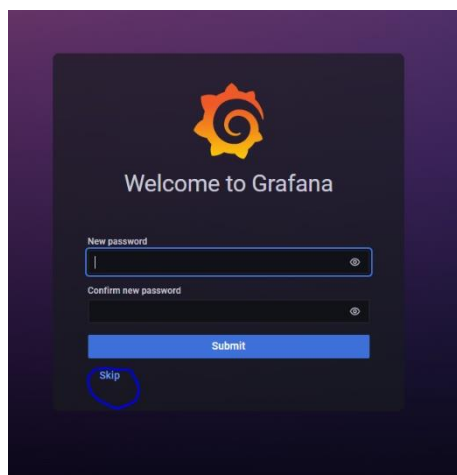
To access this dashboard, we could simply use a web browser with localhost using the following URL.

<http://localhost:3000>

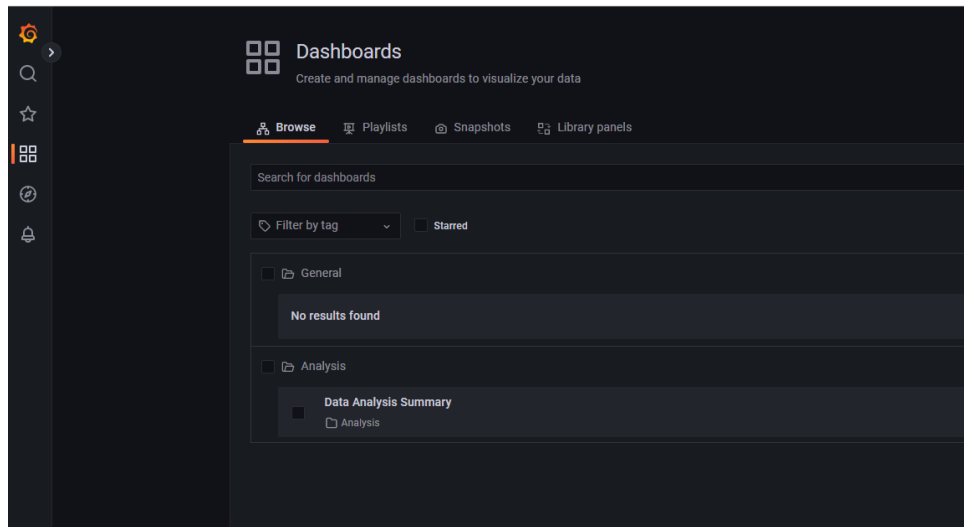
then we will get a UI like this



Default username password would be **admin:admin**



Then skip the resetting password, and next we can browse the dashboard as follows.



## References

<https://masum-math8065.medium.com/how-to-convert-categorical-data-into-numeric-in-pyspark-2202407f5fac>

<https://blog.devgenius.io/apache-storm-on-kubernetes-b0cda6b39546>

<https://www.elastic.co/guide/en/elasticsearch/hadoop/current/mapreduce.html>

<https://cazton.com/consulting/enterprise/lambda-architecture>

<https://www.elastic.co/blog/building-real-time-dashboard-applications-with-apache-flink-elasticsearch-and-kibana>

<https://www.elastic.co/what-is/elasticsearch>